Git and GitHub: A Guide for Economists

Frank Pinter

22 February 2019

Outline

The importance of version control

Using version control

Using Git for collaboration

What is version control?

Version control is a way to keep track of changes to code, text, and documents. And data and outputs.

▶ It gives you an organized revision history

▶ It lets you experiment without fear

▶ It lets you go back and forth between many different versions of the same file, and see a list of the differences

▶ It makes (the technical aspects of) collaboration a breeze

▶ It lets you and your collaborators work on different versions and then merge them

What is Git?

- ▶ Git is a program that does version control
- ▶ It is the most popular version control program in software development
- ▶ It is easy to set up and get started
- ▶ There are many programs that add intuitive interfaces on top of Git
- ▶ Git integrates seamlessly with online collaboration tools like GitHub and GitLab

Table of Contents

Using it yourself

Git isn't just useful for collaboration. It also helps you keep your own projects organized.

Life without version control

You're writing a paper and you have a regression.

- ▶ Advisor 1 tells you to include a certain variable. You put in lots of work to get the data, clean it, merge it, change the specification, and re-run.

- ▶ Advisor 2 tells you that variable is dumb. You remove it.

- ▶ Then Advisor 2 changes their mind.

What do you do?

Life without version control

Do you keep every specification you ever tried?

▶ The code and the outputs?

▶ What if you discover a coding error that affects many of your specifications?

▶ How do you organize all the files?

Version control 0.1: putting dates on things

Does this look familiar?

`run_regs_11_17_2018_v4_final_final.do`

"Not one piece of commercial software you have on your PC, your phone, your tablet, your car, or any other modern computing device was written with the 'date and initial' method." (Gentzkow and Shapiro)

Version control 0.2: Dropbox

- ▶ Dropbox keeps a crude version history.
  - ▶ But there are no labels or comments, and it's not easy to see the differences between files.
  - ▶ So if you want to dig up "the version where I had that other variable" you have to manually look through a bunch of versions.
  - ▶ And good luck if you changed two scripts, not just one.
- ▶ Dropbox lets you and your collaborators stay in sync.
  - ▶ What if you and your coauthor try to change the same script at the same time?
  - ▶ What if you are trying one change and, at the same time, your coauthor is trying a different change?

Version control 0.2: Dropbox

A Post It note spotted on a grad student's desk:

> *Don't forget! At 10:18 am on November 17th, we changed the specification to add new variable.*

Don't live this way.

Table of Contents

Why use Git?

Git is the dominant version control system today. There are others, but they're generally more work with no benefit.

Getting started

1. Install Git (Linux, Mac, Windows)
2. Git comes with a command line interface (powerful!). You might want to add a graphical interface to make things easier:

   - GitKraken
     - The screenshots in this presentation are from GitKraken
   - GitHub Desktop
   - RStudio (for R projects)

The Git model

1. You do work in your **working directory**.
2. Then you add it to your **staging area**.
3. Once you've staged all your changes for one discrete task, **commit** a snapshot of the staging area.
4. If you have a remote repository, **push** your commit to the remote repository.

Stage changes
git add

Commit changes
git commit

Push
git push

Unstaged changes in working directory

Staging area

Local

Remote

Commits: saving a snapshot

What is "one discrete task"? A collection of changes, across multiple files, that does *one thing*. Examples:

- ▶ Change the formatting of a variable from string to numeric, and treat it properly across multiple scripts
- ▶ Change your regression specification in code, in the output, and in your paper and supporting documentation
- ▶ Add a new function, and tests for that function

These form one **commit**, which you annotate with a detailed **commit message**. Examples:

- ▶ "Change the formatting of start date variable from string to Stata date format"
- ▶ "Add year dummies to regression specification"

The more detail, the more your future self will thank you.

Commits

Run tests before you commit

Your code should run properly when you commit.

- ▶ No runtime errors
  - ▶ Test this by running all code that changed, and everything that depends on it
  - ▶ Makefiles automate this process
  - ▶ Only skip if you are sure you didn't change anything important
- ▶ No compilation errors (including LaTeX)
- ▶ Tests should pass
- ▶ Output should be consistent with what you've written
  - ▶ Don't report a negative regression coefficient, and write in words that the estimated coefficient is positive

But it's better to have frequent commits (that might have small mistakes) than to have giant, infrequent commits.

You can also do partial commits (`git add --patch`) which are somewhat advanced.

What should I include?

1. At a minimum:
   - Code (`.do`, `.R`, `.m`, `.jl`, and so on)
   - Text files (`.txt`)
   - LaTeX documents (`.tex`)
2. I also recommend:
   - Raw `.csv` datasets, if small ($<$10 MB)
3. These are binary files, so you can't see differences between versions. I recommend including them anyway.
   - PDF files
   - Word, Excel, PowerPoint files
4. Some people also include all datasets.
   - Note that GitHub doesn't allow files larger than 100 MB, or projects with total size larger than 1 GB.

For datasets, look into Git Large File Storage.

What should I exclude?

In order to avoid driving your coauthors crazy, you **must** tell Git to ignore the junk files using a file called *.gitignore*. It looks like this:
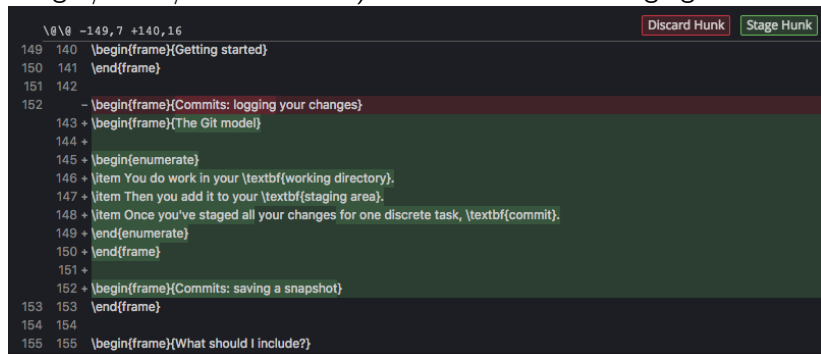
```
# Junk created by LaTeX
*.synctex.gz
*.out
*.log
# Junk created by R
.RData
# Junk created by Python
*.pyc
```

Best practice: use *.gitignore* to explicitly exclude *everything* that you don't want to include, and commit *.gitignore* like any other regular file.

GitHub maintains a list of standard *.gitignore* files for many common languages.

Viewing changes: diff

Git easily lets you see what changed in LaTeX, code (not images/PDFs/most datasets). Review this when staging!


```
\@\@ -149,7 +140,16                                          Discard Hunk   Stage Hunk
149  140   \begin{frame}{Getting started}
150  141   \end{frame}
151  142
152      -  \begin{frame}{Commits: logging your changes}
     143 +  \begin{frame}{The Git model}
     144 +
     145 +  \begin{enumerate}
     146 +  \item You do work in your \textbf{working directory}.
     147 +  \item Then you add it to your \textbf{staging area}.
     148 +  \item Once you've staged all your changes for one discrete task, \textbf{commit}.
     149 +  \end{enumerate}
     150 +  \end{frame}
     151 +
     152 +  \begin{frame}{Commits: saving a snapshot}
153  153   \end{frame}
154  154
155  155   \begin{frame}{What should I include?}
```
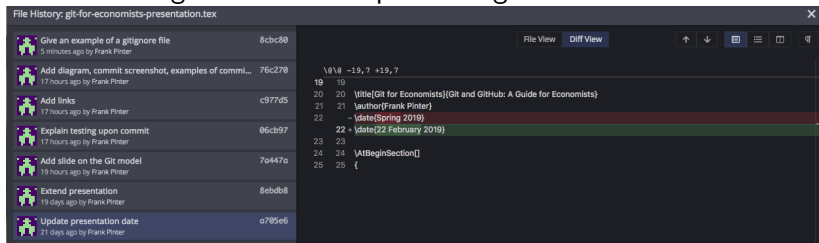
(This is the GitKraken interface, but it looks similar in any other interface)

Viewing changes: diff

You can also go back and see past changes.

Branches: trying things out

Branches are the most powerful part of Git

- ▶ By default, all the work you do goes into the "master" branch
- ▶ Want to experiment? Start a new branch
- ▶ You can switch between branches, and make commits to either branch
  - ▶ There is a catch: if you don't include your intermediate/final datasets in Git, you may need to re-make them when you switch
- ▶ If your experiment works out, commit and merge back into the master branch
  - ▶ If there are conflicts between the commits you've made on the two branches, Git will ask you to resolve them
  - ▶ This is easiest with a graphical interface like GitKraken
- ▶ If your experiment doesn't work out, delete the new branch painlessly
- ▶ The best way to learn how to use branches: practice

Keeping it local vs. using a remote repository

Git doesn't require a remote repository. You can run it 100% on your computer, with no connection to an outside server.

- ▶ This is useful if you have restrictions on your code (for example, you work with confidential health data)
  - ▶ Ask me if you have questions about using Git this way on the NBER cluster
- ▶ But a remote repository helps you keep things backed up seamlessly, and lets you collaborate with others
- ▶ You can push all your branches to the remote repository, or only some of them

Table of Contents

Repos

What's GitHub?

Brief detour: hosting services

- With a free GitHub account, you can create
  - as many private repositories as you want
  - but each private repository can only have three collaborators.
  - You can also create as many public repositories as you want (and each can have as many collaborators as you want).
- GitLab is a competing service. With a free GitLab account, you can create as many private repositories as you want, with as many collaborators as you want.
- It's easy to use GitHub for one project, and GitLab for another

Conclusion

Further reading

- ▶ Matt Gentzkow and Jesse Shapiro, "Code and Data for the Social Sciences: A Practitioner's Guide" (`https://web.stanford.edu/~gentzkow/research/CodeAndData.pdf`). See Chapter 3 for more on why you should use version control.

- ▶ Jesús Fernández-Villaverde's notes on high-performance computing (see also his class Computational Economics). Chapter 5 (`https://www.sas.upenn.edu/~jesusfv/Chapter_HPC_5_Git.pdf`) is an extended Git tutorial using the command line interface.

- ▶ Hadley Wickham's book on writing R packages. The chapter on Git and GitHub (`http://r-pkgs.had.co.nz/git.html`) is well-written and not specific to R.